

People-Intelligence Tools — Cheat Sheet

Tools 18–22 · ecos-crm-mcp · Deployed 2026-04-22

Tool Map

#	Tool	One-liner
18	<code>`add_person_observation`</code>	Write one atomic analytical insight about a contact
19	<code>`get_person_observations`</code>	Read observations, optionally filtered by type
20	<code>`compile_person_snapshot`</code>	Write versioned compiled person card (atomic)
21	<code>`get_person_card`</code>	Full composite card — primary pre-meeting entry point
22	<code>`search_brain_for_contact`</code>	Semantic BRAIN search by name at threshold 0.30

Tool 18 — add_person_observation

Required:

- `contact_id` (uuid)
- `observation_type` — `fact` | `observation` | `interpretation` | `hypothesis` | `strategy`
- `content` — standalone statement; must make sense without context
- `confidence` — integer 1–5

Optional:

- `domain_context` — `"tango"`, `"ttc"`, `"business"`, etc.
- `observed_at` — ISO timestamp; defaults to now
- `source` — defaults to `"claude-code"`
- `linked_thought_id` — BRAIN thought UUID (provenance link)

Observation Type Guide

Type	Use when	Example
<code>`fact`</code>	Verifiable, objective	"Is a licensed massage therapist"
<code>`observation`</code>	Behavioral pattern seen, not yet interpreted	"Responds to written messages faster than calls"
<code>`interpretation`</code>	Your read of what the behavior means	"Uses reconciliation framing to reopen closed channels"
<code>`hypothesis`</code>	Plausible explanation, not yet confirmed	"May be responding to abandonment panic rather than stress"
<code>`strategy`</code>	What to do in response	"Keep all contact in writing; do not meet in person"

Atomicity rule: If content contains "and" joining two distinct claims → split into two calls.

Supercession: When a new interpretation updates a prior one, note it explicitly in content: "Supersedes earlier reading that X — current read is Y."

Tool 19 — get_person_observations

Param	Required	Notes
`contact_id`	yes	
`observation_type`	no	Filter to one type
`limit`	no	Default 20

Returns observations grouped by type, newest first within each group.

Tool 20 — compile_person_snapshot

Required: *contact_id*, *snapshot_content* (readable prose)

Optional: *domains_covered* (string[]), *source_observation_ids* (uuid[]), *source_thought_ids* (uuid[])

Uses *compile_snapshot_tx* RPC — flips previous snapshot *is_current* → *false* atomically.

Never direct-insert into *person_snapshots*. Mirrors *create_billing_entry_tx* pattern.

When to compile: after seeding 5+ observations, or before a high-stakes interaction.

Tool 21 — get_person_card

Required: *contact_id* only.

Returns:

1. Contact header (name, company, role, admin status)
2. Latest compiled snapshot (if exists)
3. Recent observations grouped by type (up to 5 per type)
4. Linked BRAIN thoughts

Primary pre-meeting entry point. Run this first.

Tool 22 — search_brain_for_contact

Param	Required	Default	Notes
`contact_name`	yes	—	First name works; full name better
`domain_context`	no	—	Narrows query
`limit`	no	10	
`threshold`	no	**0.30**	Do not raise; lower threshold is intentional

Constructs query internally: *[name] [domain?] person notes observations pattern behavior*

Why 0.30: Person/TASTE/framework entries use analytical language — they don't surface reliably at the standard 0.38 threshold. Validated decision; do not adjust.

Workflow Patterns

Seed a contact from a memory extraction

1. Run extraction prompt in claude.ai (see *people-intel-prompts.md*)
2. Paste structured output into Claude Code session
3. *search_contacts* → get *contact_id*
4. *add_person_observation* per atomic item (CURRENT: yes only)
5. *compile_person_snapshot* after 5+ observations loaded

Pre-meeting prep

1. *get_person_card* → full context
2. *search_brain_for_contact* (domain = meeting topic) → BRAIN enrichment
3. Compile updated snapshot if observations have grown since last snapshot

Post-conversation capture

1. *log_interaction* → the event
2. *add_person_observation* → new patterns noticed (type: observation or interpretation)
3. If a pattern changes a prior interpretation: new *interpretation* entry noting what it supersedes

BRAIN dedup / staleness review

1. *search_brain_for_contact* at threshold 0.30
2. Compare results against *get_person_observations*
3. Flag duplicate BRAIN entries and entries superseded by newer observations
4. Update or delete via *mcp__open-brain__update_thought* / *mcp__open-brain__delete_thought*